

Interaction filtering - A novel approach to social recommendation

Riley Kidd

A subthesis submitted in partial fulfillment of the degree of
Bachelor of Software Engineering at
The Department of Computer Science
Australian National University

October 2012

© Riley Kidd

Typeset in Palatino by \TeX and $\text{\LaTeX} 2_{\epsilon}$.

Except where otherwise indicated, this thesis is my own original work.

Riley Kidd
11 October 2012

Abstract

Social networks provide a wide array of user specific interactions, profile information and user preferences. This thesis attempts to decipher which user interactions or preferences are truly indicative of 'likes', this information is then leveraged to allow for binary classification of user specific links with the goal of discovering the ideal combination of traits for prediction.

The success of our predictions are evaluated using a number of machine learning algorithms including, *Naive Bayes*, *Logistic Regression* and *Support Vector Machines*, results are also compared to previous work using *Matchboxing* and *Social Matchboxing* techniques. The data set is sourced from a set of over 100 Facebook users and their interactions with over 30,000 friends during a four month period.

Our analysis has shown that

Contents

Abstract	v
1 Introduction	1
1.1 Objectives	1
1.2 Contributions	1
1.3 Outline	2
2 Background	3
2.1 Facebook	3
2.2 LinkR	3
2.3 Notation	5
2.4 Feature Sets	5
2.5 Previous Work	5
2.5.1 Content Based Filtering	6
2.6 Classification Algorithms	6
2.6.1 Constant	6
2.6.2 Social Match Box	6
2.6.3 Naive Bayes	6
2.6.4 Logistic Regression	6
2.6.5 Support Vector Machine	6
2.7 Training and Testing	7
2.8 Evaluation Metrics	7
3 Interactions	9
3.1 User Interactions	9
3.2 Conversation	12
3.2.1 Outgoing	14
3.2.2 Incoming	15
3.3 Conclusion	15
4 User Preferences	21
4.1 Introduction	21
4.2 Demographics	21
4.3 Traits	26
4.4 Groups	31
4.5 Pages	35

5	Model Combinations	39
5.1	Introduction	40
5.2	Positive Feature Selection	40
5.3	Bayesian Model Averaging	40
5.4	Summary	40
6	Conclusions	41
6.1	Future Work	41
	Bibliography	43

Introduction

An individual's social presence on the web is continually expanding, with the emergence of services such as Facebook, Myspace, LinkedIn, Twitter and Google+ what defines a user and their online social interactions (messages, posting, commenting, etc) and preferences (demographics, group memberships, likes, etc) is an ever expanding graph structure of verbose social content. Providing a myriad of expanding social information and user driven content.

Our question becomes, how can we exploit this information to decipher which user interactions or preferences are most indicative of user likes? Can we leverage this information and generalise over a social population?

1.1 Objectives

The primary objective of this thesis is to contrast and compare differing user feature sets. Using the machine learning concepts of *Naïve Bayes*, *Logistic Regression* and *Support Vector Machines* compared with our appropriate baselines of *Social Matchbox* and *Constant Classifiers*.

Using these classification techniques we will compare both user interaction features *interactions*, *incoming/outgoing messages* and user preference features *demographics*, *traits*, *groups*, *pages*. To analyse the effect each feature set has on classification.

These algorithms will also be tested against a subset of the data, using a hold out technique to only test against items which has been liked by some friend.

Finally, we will analyse the effect of combining successful user feature sets together and the results of combining successful models together using a *Bayesian Model Averaging* approach.

1.2 Contributions

In the preceding section, we outlined different user interaction and preference features. Our specific contributions made during this thesis show

- Both *interactions* and *incoming/outgoing messages* are not more predictive than previously used *Social Matchboxing* techniques.

- Each user preference of *demographics, traits, groups, pages* contributed to a better result than our baselines.
- Combining user preferences with a hold out technique for user likes results in a significant improvement from our baselines.
- Combination of user preferences is better
- Model combination can encapsulate these ideas

Overall, we provide a methodology which improves upon previous work and offers a way to combine positively contributing aspects of different feature sets in our data.

1.3 Outline

The remaining chapters in this thesis are organised as follows:

- **Chapter 2:** We first outline appropriate background information for the reader. Including information pertaining to the source of the data set, notation used throughout this thesis, previous work in this area and our research approach and methodology
- **Chapter 3:** In this chapter we discuss different feature sets for user interactions and the results from applying these feature sets in comparison with our baselines.
- **Chapter 4:** Similarly as above, however we discuss user preferences
- **Chapter 5:** In this chapter we discuss results from combining different feature sets and models
- **Chapter 6:** Finally, we draw the work done throughout this thesis to a conclusion and offer avenues for future work in this area.

All chapters combined, this thesis represents a novel approach to which feature sets are predictive of user likes and offer an approach to combining positive components into a useful classifier.

Background

In the following, we define the source of our data set, notation used throughout this thesis, our choice of prediction algorithms and our testing methodology.

2.1 Facebook

Facebook is the largest and most active social media service in the world. Facebook users can create a profile containing personal preferences and have friendships and interactions between other users. These interactions can also be liked or commented on by other users.

2.2 LinkR

NICTA developed a Facebook app named LinkR¹ which would make recommendations to app users and record whether or not the user liked the item.

The dataset includes information about each app user as well as a subset of available data about their friends.

The LinkR Facebook app was used to collect information about users, their interactions and preferences. The data set contains information about app users as well as a sub-set of visible information about their friends. The app tracked and stored information for over 100 app users and their 39,000+ friends.

The four main interactions between users are posts (posting an element on a friends' wall), tags (being mentioned in a friends post or comment), comments (written data on a post) and likes (clicking a like button if a user likes a post or comment). The table below outlines data collected during app trials.

The table below summarises the data collected from both app users and their friends.

Pertanent user data which is collected by the LinkR app includes:

- Gender

¹The main developer of the LinkR Facebook App is Khoi-Nguyen Tran, a PhD student at the Australian National University. Khoi-Nguyen wrote the user interface and database crawling code for LinkR.

App Users	Posts	Tags	Comments	Likes
Wall	27,955	5,256	15,121	11,033
Link	3,974	-	5,757	4,279
Photo	4,147	22,633	8,677	5,938
Video	211	2,105	1,687	710
App Users and Friends	Posts	Tags	Comments	Likes
Wall	3,384,740	912,687	2,152,321	1,555,225
Link	514,475	-	693,930	666,631
Photo	1,098,679	8,407,822	2,978,635	1,960,138
Video	56,241	858,054	463,401	308,763

Table 2.1: Data records for interactions between users. Rows are the type of interaction, columns are the context.

- Age
- Hometown
- Locale
- Group Memberships
- Page Likes
- Favourite Activities
- Favourite Books
- Favourite Athletes
- Favourite Teams
- Inspirational People
- Interests
- Favourite Movies
- Favourite Music
- Favourite Sports
- Favourite Television Shows
- School Information
- Work Information
- Messages data

2.3 Notation

The mathematical notation used by our classifiers during this thesis are outlined below.

- N users.
- M items.
- User features F of size i user features.
- A dataset D comprised of $D = \{(n, m, f_i) \rightarrow y\}$ with the binary response $y \in \{0, 1\}$ where 0 represents a dislike and 1 represents a like.

2.4 Feature Sets

The feature sets in x can be any of the following, which are discussed further in :::

- Interactions
- Demographics
- Traits
- Groups
- Pages
- Outgoing Messages
- Incoming Messages

2.5 Previous Work

While many Facebook users have a friend count which is close to the human real word limit, known as the Dunbar number [Hill and Dunbar 2003], this work shows that user interactions are focused on a much smaller subset of their friends.

[Backstrom et al. 2011] studied two types of user uses of Facebook, explicit communication interaction and viewing attention. Communication is focused on a limited subset of friends whilst viewing attention is dispersed among a much larger set. This supports the approach of testing a wide array of user interactions and preferences, as each users preferences are driven by where their attention is focused.

2.5.1 Content Based Filtering

Content based filtering (CBF) [?] is an extension of the technique of Collaborative filtering (CF) [?] which predicts whether a user will like an item via information about that users' preferences as well as that of other users, CBF extends this approach by generalising from the item features which the user has explicitly liked or disliked.

Baesd on previous user trials [?] Social Matchbox was the best performing algorithm in live user trials, gaining more likes then dislikes.

2.6 Classification Algorithms

This analysis makes use of the results from a number of different classification algorithms which are outlined below.

2.6.1 Constant

The constant predictor returns a constant result irrespective of the feature vectors selected from above. The most common result in our data set is *False* and hence the *False* predictor is displayed in our analysis.

2.6.2 Social Match Box

2.6.3 Naive Bayes

Naive Bayes (NB) is a basic predictor which involves applying Bayes' theorem using independence assumptions between each feature in x .

The NB implementation used during this thesis is an implementation previously devised by *Scott Sanner* [?].

2.6.4 Logistic Regression

Logistic Regression (LR) predicts the odds of being either a like or a dislike by converting a dependent variable and one or more continuous independent variable(s) into probabability odds.

The LR implementation used during this thesis is *LingPipe* [?].

2.6.5 Support Vector Machine

The *Support Vector Machine* (SVM) is a supervised learning machine based on a set of basis functions which help construct a separating hyperplane between the data points. Training involves buidling the relevant hyperplanes which can then be used for testing. Each data point is classified depending on which side of the hyperplane it falls.

The SVM implementation used during this thesis is *SVMLibLinear* [Chang and Lin 2011].

2.7 Training and Testing

All evaluation is done using 10 fold cross validation wherein the data is partitioned into 10 complimentary subsets, each subset is composed of two separate parts one section is used for training (80%) and the other (20%) is used for testing. This is performed on 10 distinct subsets and the results are averaged across each fold.

2.8 Evaluation Metrics

When evaluating the success of each method at correctly predicting the classification, the following metrics will be used.

- A *true positive* prediction refers to when the classifier correctly identifies the class as true.
- A *false positive* occurs when the prediction is true, but the true class was false.
- A *false negative* occurs when the prediction is false but the actual class is true.

Accuracy relates to the closeness to the true value. In the context of our results, the accuracy refers to the number of correct classifications divided by the size of the data set.

$$\text{accuracy} = \frac{\text{number of correct classifications}}{\text{size of the test data set}}$$

Precision relates to the number of retrieved predictions which are relevant. In the context of our results, the precision refers to the number of true positive predictions divided by the sum of the true positive and false positive predictions.

$$\text{precision} = \frac{\text{number of true positives}}{\text{number of true positives} + \text{number of false positives}}$$

Recall refers to the number of relevant predictions that are retrieved. In the context of our results, recall refers to the number of true positive predictions divided by the sum of the true positive and false negative predictions.

$$\text{recall} = \frac{\text{number of true positives}}{\text{number of true positives} + \text{number of false negatives}}$$

The f-score combines and balances both precision and recall and is referred to as the weighted average of both precision and recall.

$$\text{f-score} = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

The main metric we use for analysis in this thesis is accuracy.

Interactions

The user interactions we examine in this thesis can be broken down into two distinct groups, interactions between users and messages sent between users.

3.1 User Interactions

There are a number of potential interactions between users under the Facebook paradigm.

- Direction: Incoming or outgoing Given the approach of [Saez-Trumper et al. 2011] where interaction directionality was shown to be highly reflective of user preferences we decipher between incoming and outgoing interactions
- Modality: Links, posts, photos or videos
- Type: Comment, tag or like

In this case, our feature vector x_i is comprised of the crossproduct, where:

$$i = \{incoming, outgoing\} \times \{post, photo, video, link\} \times \{comment, tag, like\}$$

The alters of i can then be defined as all users who have interacted with the current user via some interaction i . The column is set to 1 if any of the alters defined by the current set i have also liked the item associated with the user, otherwise it is set to 0.

Applying our classification algorithms defined above and this feature set we obtain:

Compared with our baselines, user interactions do not appear to help our classification.

However, if we use only a subset of the data, holding out on only elements which have at least k likes among the group of alters we can compare using these exposure curves [Romero et al. 2011] which can provide some indicator of a users exposure to a liked item.

Having one user liking an item is simply enough.




Figure 3.1: Predictors paradigm




Figure 3.2: Accuracy results using the *user interactions* feature set




Figure 3.3: Accuracy results for an exposure curve using the *user interactions* feature set

3.2 Conversation

Given the nature of Facebook, it is possible for users to post messages to other users.

These messages can be broken down based on their directionality, either *outgoing* which are words sent to other users or *incoming* which are words received from other users.

The most commonly used words are below.

Rank	Word	Frequency
1	:)	292,733
2	like	198,289
3	good	164,387
4	thanks	159,238
5	one	156,696
6	love	139,939
7	:p	121,904
8	time	106,995
9	think	106,459
10	see	103,690
11	nice	99,672
12	now	94,947
13	well	92,735
14	happy	84,381
15	:d	83,698
16	much	78,719
17	oh	77,321
18	yeah	76,564
19	back	76,032
20	great	70,514
21	going	70,447
22	still	68,245
23	new	67,430
24	day	65,579
25	come	63,837
26	;)	62,936
27	year	61,771
28	look	60,608
29	yes	59,774
30	want	59,514
31	tag	58,633
32	hahaha	57,448
33	also	56,414
34	need	55,921
35	make	54,949
36	sure	54,395
37	thank	54,112
38	people	53,211
39	miss	53,182
40	guys	52,855
41	right	52,112
42	best	51,941
43	awesome	51,663
44	hope	50,980
45	2	50,720
46	next	50,375
47	work	49,459
48	way	49,358
49	man	49,101
50	:(48,184
51	j3	47,985
52	even	47,480
53	4	46,068
54	us	45,919
55	pretty	44,804
56	hey	44,614
57	say	44,315
58	better	43,357
59	thanx	42,639
60	bro	41,187
61	take	41,081
62	always	40,457
63	wow	40,452
64	pic	40,185
65	though	40,032
66	actually	39,565
67	last	39,175
68	thats	38,833
69	cool	37,844
70	dear	37,328
71	ok	36,441
72	sorry	36,345
73	never	36,000
74	thing	35,941
75	first	35,785
76	looks	35,496
77	night	35,475
78	thought	34,458
79	photo	33,989
80	&	33,902

Table 3.1: Top conversation content data for all users

3.2.1 Outgoing

First we need to figure out the number of outgoing messages words are optimal for our classifiers. Given the size of potential messages we decided to test within a range of (100-1000) where each word is based on its top frequency.




Figure 3.4: Accuracy results for different *outgoing words* sizes

The most predictive top outgoing words sizes for our classifiers are:

- Naive Bayes: 500
- Logistic Regression: 200
- Support Vector Machine: 900

Testing our classifiers using their optimal x_i based on the above analysis we obtain: These results do not show an improvement over our baselines.

However when using an exposure curve we find that having some user liking our item is enough to improve classification.




Figure 3.5: Accuracy results using the *incoming words* feature set

3.2.2 Incoming

Similarly for incoming messages we need to discover which is the best.

- Naive Bayes: 300
- Logistic Regression: 100
- Support Vector Machine: 1000

Testing using these optimal message size groups we obtain.

Again incoming messages themselves are not predictive.

3.3 Conclusion

User interactions in themselves are not predictive of user likes.

[?] concluded that it is less important what users say, then who they interact with.

This is supported by [?] who found that virtual interactions help reveal common interests, while real world interactions helps to support friendships.




Figure 3.6: Accuracy results for an exposure curve using the *outgoing words* feature set




Figure 3.7: Accuracy results for different *incoming words* sizes




Figure 3.8: Accuracy results using the *incoming words* feature set




Figure 3.9: Accuracy results for an exposure curve using the *incoming words* feature set

User Preferences

In this section we will discuss the effects of using user preferences as the feature set.

4.1 Introduction

Facebook allows users to provide a vast array of personal traits and interests on their Facebook page.

Including:

- Demographics - age, gender, location, etc
- Group Memberships
- Personal Preferences - favourite books, favourite athletes, favourite sports, inspirational people, personal interests, etc
- Conversation Data - words sent, words received

In this section we will try to uncover which user preferences are indicative of item likes.

4.2 Demographics

Gender breakdown in the data set:

Male	Female	Undisclosed
85	33	1

Table 4.1: Gender breakdown

There is a clear male bias in the data set.

[Backstrom et al. 2011] have shown that different genders have differing tendencies to disperse interactions across genders.

Birthday breakdown in the data set:

Birthdays are grouped in a distinct range, most users in this data set are in the age range of 18 – 30.

Year	Frequency
Undisclosed	1
1901-1905	1
1906-1910	0
1911-1915	1
1916-1920	0
1921-1925	0
1926-1930	0
1931-1935	0
1936-1940	1
1941-1945	0
1946-1950	0
1951-1955	0
1956-1960	2
1961-1965	1
1966-1970	4
1971-1975	10
1976-1980	12
1981-1985	25
1986-1990	34
1991-1995	25
1996-2000	2

Table 4.2: Birthday breakdown

Undisclosed and anomlies (100 years plus) something.

Location breakdown in the data set:

Location	Frequency
Undisclosed	33
Ahmedabad, India	1
Bangi, Malaysia	1
Bathurst, New South Wales	1
Bellevue, Washington	1
Braddon, Australian Capital Territory, Australia	1
Brisbane, Queensland, Australia	2
Canberra, Australian Capital Territory	56
Culver City, California	1
Frederick, Maryland	3
Geelong, Victoria	1

Table 4.3: Location breakdown

Given the fact that most users are either situated in the ACT (location of the app development and deployment) or are undisclosed, location information in this data set will not be used.

The feature vector used for demographics is comprised of:

- Whether the user is male
- Whether the user is female
- Whether the user and any user in the alters set share the same gender
- Whether the user and any user in the alters set share the same birthrange

The demographics feature set shows positive results, which are nearly comparable to the baseline of SMB.

The exposure curve shows a positive correlation between friend likes.

4.3 Traits

The traits feature set includes:

- Activities
- Books
- Athletes
- Teams
- Inspirational People

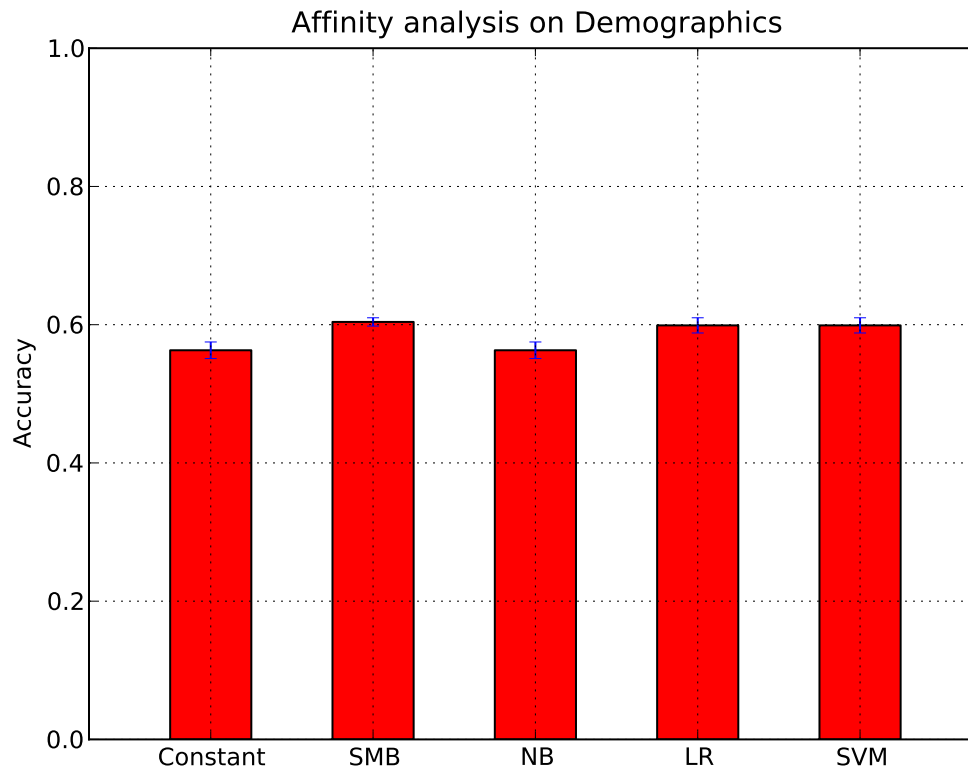


Figure 4.1: Accuracy results using the *demographics* feature set

- Interests
- Movies
- Sports
- Television
- School Relationships
- Work Relationships

The traits feature set shows an improvement over our SMB baseline in the LR and SVM case.

This tren continues through exposure curve.




Figure 4.2: Accuracy results for an exposure curve using the *demographics* feature set

Count	Activity
10	Sleeping
5	Eating
5	Reading
4	Running
4	Cycling
4	Minecraft
4	Programming
3	Android
3	Cooking
3	Video Games
3	Xbox 360
3	Piano
3	Guitar
3	Badminton
3	Chocolate

Table 4.4: Top activities for app users

Count	Inspirational Person
2	Alan Turing
1	Bender
1	Maurice Moss
1	Steve Jobs
1	Sean Parker
1	Pope Benedict XVI
1	Martin Luther
1	Alistair McGrath
1	St Augustine
1	Dennis Ritchie
1	Linus Torvalds
1	Richard Stallman
1	C. S. Lewis
1	Mike Oldfield
1	Ryan Giggs

Table 4.5: Top inspirational people

Count	Book
7	Harry Potter
4	The Bible
3	Harry Potter series
3	Discworld
3	That's 3 minutes of solid study, think I've earned 2hrs of Faceboook time
3	Freakonomics
3	Tomorrow when the War Began
2	Magician
2	Hitchhiker's Guide To The Galaxy
2	The Discworld Series
2	Terry Pratchett
2	Terry Pratchett
2	George Orwell
2	Lord Of The Rings
2	Goosebumps

Table 4.6: Top books for app users

Count	Athlete
4	Roger Federer
4	Rafael Nadal
3	Maria Sharapova
2	Leo Messi
1	Andy Schleck
1	Chrissie Wellington
1	Emma Snowsill
1	Emma Moffatt
1	Brbara Riveros
1	The Brownlee Brothers
1	Marie Slamtoinette #1792
1	Wayne Rooney
1	"you are what you eat" " I dont remember eating a Tank."
1	Nemanja Vidic
1	Ryan Giggs

Table 4.7: Top athletes for app users

Count	Interest
5	Movies
5	Music
3	Cooking
3	Sports
2	Psychology
2	Internet
2	Video Games
2	Martial arts
2	Literature
2	Economics
2	Tennis
2	Badminton
2	Artificial intelligence
2	Computers
2	Travel

Table 4.8: Top interests for app users

Count	Music
9	Daft Punk
9	Muse
8	Michael Jackson
8	Pink Floyd
8	Lady Gaga
7	Linkin Park
7	Avril Lavigne
6	Radiohead
6	Rihanna
6	Coldplay
6	Green Day
6	Katy Perry
6	Taylor Swift
5	Gorillaz
5	Queen

Table 4.9: Top music for app users

Count	Movie
9	Inception
8	Avatar
8	Fight Club
7	The Lord of the Rings Trilogy (Official Page)
6	Star Wars
6	I wouldnt steal a car, But i'd download one if i could
6	WALL-E
6	Scott Pilgrim vs. the World
6	Toy Story
6	Shrek
5	Batman: The Dark Knight
5	Harry Potter
4	The Matrix
4	The Social Network Movie
4	Monsters, Inc.

Table 4.10: Top movies for app users

Count	Sport
8	Badminton
5	Basketball
3	Cycling
3	Volleyball
2	Starcraft II
2	Football en salle
2	Swimming
2	Towel Baseball
2	Tennis
1	Soccer
1	Taekwondo
1	Rock climbing
1	In The Groove
1	Darts
1	Table tennis

Table 4.11: Top sports for app users

Count	Television Show
20	The Big Bang Theory
19	How I Met Your Mother
14	The Simpsons
13	Top Gear
12	Futurama
12	Scrubs
11	Black Books
10	Black Books
10	South Park
10	Family Guy
9	The Daily Show
8	The IT Crowd
8	FRIENDS (TV Show)
7	True Blood
7	MythBusters

Table 4.12: Top television shows for app users

Count	Team
5	Manchester United
2	Bear Grylls cameraman appreciation society
2	Real Madrid C.F.
2	Liverpool FC
1	Leopard Trek
1	British Triathlon
1	TeamCWUK
1	Surly Griffins
1	Canberra Raiders
1	Kolkata Knight Riders
1	Brisbane Roar FC
1	Brisbane Broncos
1	Cricket Australia
1	— Manchester United Fans —
1	Juventus

Table 4.13: Top teams for app users




Figure 4.3: Accuracy results using the *user traits* feature set




Figure 4.4: Accuracy results for an exposure curve using the *user traits* feature set

4.4 Groups

The most popular groups for app users are shown below.

Group Name	Frequency
27	ANU StalkerSpace
20	Facebook Developers
15	ANU CSSA
14	CSSA
13	Australian National University
11	ANU - ML and AI Stanford Course
10	iDiscount ANU
10	Our Hero: Clem Baker-Finch
9	Students In Canberra
7	I grew up in Australia in the 90s
7	Grow up Australia - R18+ Rating for Computer Games
7	ANU Engineering Students' Association (ANUESA) 2010
7	ANU Postgraduate and Research Student Association (PARSA)
6	No, I Don't Care If I Die At 12AM, I Refuse To Pass On Your Chain Letter.
6	No Australian Internet Censorship
6	The Chaser Appreciation Society
6	Feed a Child with a Click
6	ANU Mathematics Society
6	ANU International Student Services, CRICOS Provider Number 00120C
6	2011 New & Returning Burton & Garran Hall
5	If You Can't Differentiate Between "Your" and "You're" You Deserve To Die
5	Keep the ANU Supermarket!!!
5	If 1m people join, girlfriend will let me turn our house into a pirate ship
5	The Great Australian Internet Blackout
5	When I was your age, Pluto was a planet.

Table 4.14: App users groups breakdown for range 5+




Figure 4.5: Accuracy results for different *groups* sizes




Figure 4.6: Accuracy results using the *groups* feature set




Figure 4.7: Accuracy results for an exposure curve using the *groups* feature set

4.5 Pages

Page Name	Frequency
33	ANU Computer Science Students' Association (ANU CSSA) 2011
32	The Australian National University
31	ANU Stalkerspace
21	Humans vs Zombies @ ANU
20	The Big Bang Theory
19	Australian National University
19	How I Met Your Mother
18	ANU LinkR
18	ANU ducks
17	Australian National University Students' Association
16	Google
15	Google Chrome
15	ANU XSA
15	Facebook
14	YouTube
14	The Simpsons
13	Portal
13	Top Gear
13	Music
13	ANU Memes
12	Futurama
12	Scrubs
12	ANU O-Week 2012: Escape to the East
12	The Stig
11	Black Books

Table 4.15: App users groups breakdown for range 5+




Figure 4.8: Accuracy results for different *pages* sizes




Figure 4.9: Accuracy results using the *pages* feature set




Figure 4.10: Accuracy results for an exposure curve using the *pages* feature set

Model Combinations




Figure 5.1: Predictors paradigm




Figure 5.2: Predictors paradigm

5.1 Introduction

5.2 Positive Feature Selection

5.3 Bayesian Model Averaging

5.4 Summary

Conclusions

In conclusion this offers the most important results for the initial questions that motivated this thesis.

6.1 Future Work

- Passive likes:
- Cold start:
- General user set:

Bibliography

- BACKSTROM, L., BAKSHY, E., KLEINBERG, J., LENTO, T., AND ROSENN, I. 2011. Center of attention: How facebook users allocate attention across friends. *ICWSM'11* (2011). (pp.5, 21)
- BAKSHY, E., ROSENN, I., MARLOW, C., AND ADAMIC, L. 2012. The role of social networks in information diffusion. *Facebook report*, <http://www.scribd.com/facebook>.
- CHANG, C.-C. AND LIN, C.-J. 2011. LIBSVM: A library for support vector machines. *ACM Transactions on Intelligent Systems and Technology* 2, 27:1–27:27. (p.6)
- HILL, R. AND DUNBAR, R. 2003. Social network size in humans. *Human Nature* 14, 1, 53–72. (p.5)
- LERMAN, K. AND GHOSH, R. 2010. Information contagion: An empirical study of the spread of news on digg and twitter social networks (2010).
- MA, KING, AND LYU. 2009. Learning to recommend with social trust ensemble. In *SIGIR-09* (2009).
- MA, H., YANG, H., LYU, M. R., AND KING, I. 2008. Sorec: Social recommendation using probabilistic matrix factorization. In *CIKM-08* (2008).
- ROMERO, D. M., MEEDER, B., AND KLEINBERG, J. 2011. Differences in the mechanics of information diffusion across topics: idioms, political hashtags, and complex contagion on twitter. *WWW '11* (2011), pp. 695–704. ACM. (p.9)
- SAEZ-TRUMPER, D., NETTLETON, D., AND BAEZA-YATES, R. 2011. High correlation between incoming and outgoing activity: A distinctive property of online social networks? In *Fifth International AAAI Conference on Weblogs and Social Media*, *ICWSM '11* (2011). (p.9)
- YANG, LONG, SMOLA, SADAGOPAN, ZHENG, AND ZHA. 2011. Like like alike: Joint friendship and interest propagation in social networks. In *WWW-11* (2011).