Social Affinity Filtering: Recommendation through Fine-grained Analysis of User Interactions and Activities

Suvash Sedhain
ANU & NICTA
Canberra, Australia
ssedhain@nicta.com.au

Scott Sanner
NICTA & ANU
Canberra, Australia
ssanner@nicta.com.au

Lexing Xie
ANU & NICTA
Canberra, Australia
lexing.xie@anu.edu.au

Riley Kidd
ANU
Canberra, Australia
rileyjkidd@gmail.com

Khoi-Nguyen Tran
ANU
Canberra, Australia
kndtran@cs.anu.edu.au

Peter Christen
ANU
Canberra, Australia
peter.christen@anu.edu.au

ABSTRACT

Content recommendation in social networks poses the complex problem of learning user preferences from a rich and complex set of interactions (e.g., likes, comments and tags for posts, photos and videos) and activities (e.g., favourites, group memberships, interests). While many social collaborative filtering approaches learn from aggregate statistics over this social information, we show that only a small subset of user interactions and activities are actually useful for social recommendation, hence learning *which* of these are most informative is of critical importance. To this end, we define a novel social collaborative filtering approach termed social affinity filtering (SAF). On a preference dataset of Facebook users and their interactions with 37,000+ friends collected over a four month period, SAF learns which fine-grained interactions and activities are informative and outperforms state-of-the-art (social) collaborative filtering methods by over 6% in prediction accuracy; SAF also exhibits strong cold-start performance. In addition, we analyse various aspects of fine-grained social features and show (among many insights) that interactions on video content are more informative than other modalities (e.g., photos), the most informative activity groups tend to have small memberships, and features corresponding to "long-tailed" content (e.g., music and books) can be much more predictive than those with fewer choices (e.g., interests and sports). In summary, this work demonstrates the substantial predictive power of fine-grained social features and the novel method of SAF to leverage them for state-of-the-art social recommendation.

Categories and Subject Descriptors

H.3.3 [**Information Search and Retrieval**]: Information Filtering

Keywords

social networks, collaborative filtering, recommender systems

1. INTRODUCTION

Online social networks such as Facebook record a rich set of user preferences (likes of links, posts, photos, videos), user traits, interactions and activities (conversation streams, tagging, group memberships, interests, personal history, and demographic data). This presents myriad new dimensions to the recommendation problem by making available a rich labeled graph structure of social interactions and content from which user preferences can be learned and new recommendations can be made.

Most existing recommendation methods for social networks aggregate this rich social information into a simple measure of user-to-user interaction [9, 19, 22, 23, 34, 20, 21]. But in aggregating all of these interactions and common activities into a *single* strength of interaction, we ask whether important preference information has been discarded? Indeed, the point of departure for this work is the hypothesis that different fine-grained interactions (e.g. commenting on a wall or getting tagged in a video) and activities (e.g., being a member of a university alumni group or a fan of a TV series) *do* represent different preferential *affinities* between users, and moreover that effective *filtering* of this information (i.e., learning which of these myriad fine-grained interactions and activities are informative) will lead to improved accuracy in social recommendation.

To quantitatively validate our hypotheses and evaluate the informativeness of different fine-grained features for social recommendation, we have built a Facebook App to collect detailed user interaction and activity history available through the Facebook Graph API along with user preferences solicited by the App on a daily basis. Given this data, (1) we define a novel recommendation method called *social affinity filtering (SAF)*, where we learn to predict whether a user (ego) will like an item based on the surrogate item preferences of others (alters) who share fine-grained interactions or activities with the ego, and (2) we analyse the relative informativeness of these fine-grained interaction and activity features across a variety of dimensions.

In the four months that our App was active, we collected data for a set of Facebook app users and their full interactions with 37,000+ friends along with 22 distinct types of interaction and users activity for 3000+ groups, 4000+ favourites, and 10,000+ pages. In subsequent sections that outline our experimental methodology and results in detail, we make the following critical observations:

- **Overall performance:** We found that SAF significantly outperforms numerous state-of-the-art collaborative filtering and social recommender systems by over 6% in accuracy using just *page* (like) features.

- **Privacy vs. performance:** Because the reluctance of a user to install an App increases with the number of permissions requested, the above results suggest that an SAF-based social recommendation App need only request permissions for a user's *page* likes in order to achieve state-of-the-art recommendation accuracy.

- **Cold-start performance:** Since SAF does not condition its predictions on individual user IDs, we show that it exhibits strong *cold-start* performance for users *without* expressed item preferences as long as those users have interactions or shared activities with users who have expressed item preferences.

- **Big data scalability:** We implement SAF as a simple linear classifier that can be used in conjunction with a variety of classification methods (e.g., naive Bayes, logistic regression, SVM) and online training algorithms amenable to real-time, big data settings.

- **Interaction analysis:** Among *interactions*, we found that those on videos are more predictive than those on other content types (photos, post, link), and that outgoing interactions (performed by the ego on the alter's timeline) are more predictive than incoming ones (performed by alters on the ego's timeline), although the level of exposure of an ego to an alter's preferences is often more important than the directionality, modality, or action underlying the interaction with the alters.

- **Activity analysis:** The most predictive activity SAGs tend to have small memberships indicating that these informative activities represent highly specialised interests. We also found features corresponding to "long-tailed" dynamic content (such as music and books) can be more predictive than those with fewer choices that add little new content over time (e.g. interests or sports).

- **Importance of social data beyond friends:** We found that *groups*, *pages*, and *favourites* make for more informative SAGs than those defined by user-to-user interactions. This is likely because the former can be applied to SAGs *over the entire Facebook population* rather than just a user's friends (where the available preference data is considerably more sparse).

Figure 1: Overview of *social affinity filtering (SAF)*: A *social affinity group (SAG)* of user u (ego) consists of a set of alternate users $\{v\}$ (alters) who have a certain *interaction* or share an *activity* membership with u. SAF learns to classify whether user u will like item i based on the observed preferences of members of each SAG of user u toward item i.

- **Social activity and item popularity vs. performance:** We analyse *how many* shared activities are needed for good performance and observe that increased activity membership correlates with increased recommendation accuracy, but that excessive item popularity coupled with high activity membership hurts the discriminative power of SAF to make good recommendations.

- **Fine-grained vs. aggregate social data:** Among activity features, a small subset proved to be much more informative than the rest. This suggests the value of learning *which* fine-grained features are predictive and sheds doubt on the efficacy of existing social recommendation methods that aggregate social information between two users into a single numerical value.

Subsequent sections demonstrate these findings in detail.

2. SOCIAL AFFINITY FILTERING

As illustrated in Fig 1, the high-level objective of this work is to predict whether or not a user u (ego) will like an item i. Specifically, the Facebook App we have built for our experimentation collects explicit like and dislike feedback for links posted on Facebook (e.g., Youtube video, news or blog item, etc.) leading to the following preference data:

$$likes(u, i) := \begin{cases} true & u \text{ clicked } like \text{ for } i \\ false & u \text{ clicked } dislike \text{ for } i \\ unknown & u\text{'s preference for } i \text{ is unobserved} \end{cases}$$

From the observed data, *social affinity filtering (SAF)* learns to predict $likes(u, i)$ based on the surrogate link preferences $likes(v, i)$ of sets of other Facebook users v who have at least one interaction or activity in common with u. The details of SAF are outlined in the following subsections.

2.1 Interactions and Activities on Facebook

In the context of Facebook, we use the term *interactions* and *activities* to refer to the range of user-user and user-community actions, respectively.

Interactions describes communication between Facebook users and can be broken down into the following dimensions:

- **Modality:** (4 possibilities) User u can interact with another user v via *links, posts, photos* and *videos* that appear in either user's timeline.

- **Action type:** (3 possibilities) A user u can *comment* or *like* user v's item. He/she can also *tag* user v on an item, often indicating that user v is present when the content is created (for photo/video/post), or to explicitly raise user v's attention for a post — with one exception in Facebook that u cannot tag a link with users.

- **Directionality:** (2 possibilities) We look at *incoming* and *outgoing* interactions, i.e., if user u comments on, tags, or likes user v's item, then this is an *outgoing* interaction for u, and an *incoming* interaction for v. Although high correlation between *incoming* and *outgoing* interactions has been observed [28], whether interaction direction affects user preferences differently is still an open question we wish to answer in this work.

Overall there are 22 possible interaction types, namely the cross-product of modalities, actions and directions, minus *link-tag-{incoming, outgoing}* since links cannot be tagged.

Activities are user interactions with Facebook communities like groups, pages, and favourites defined as follows:

- **Groups** on Facebook [1] are analogous to real-world community organisations. They allow users to declare membership and support people to organise activities, to post related content, and to have recurring discussions about them. Examples of groups include *Stanford Thai* (Fig 1 bottom left), or *Harvard Debate Club*.

- **Pages** on Facebook [2] are analogous to the homepages of people, organisations and events on the world-wide-web. They are publicly visible, and users can subscribe to the updates on the page, and also engage in discussions. Example pages include *DARPA* (an organisation, Fig 1 bottom middle), or *Beyonce* (a singer).

- **Favourites** are analogous to bookmarks (on physical books or on the web browser). They are a user-created list containing various items such as Facebook apps, books, music, and many other types of items (even pages) to indicate their interest. Example favourites include *Big Bang Theory* (TV series), or *FC Barcelona* (soccer club). Fig 1 bottom right shows a Facebook screenshot when a user adds a favourite. [3]

Our evaluation includes 3000+ *group*, 4000+ *page* and 10000+ *favourite* features as detailed in Sec 3.1.

2.2 Social Affinity Groups (SAGs)

With *interactions* and *activities* now defined, we proceed to define two types of *social affinity groups (SAGs)* of a user u that will be used as proxies for u's preferences:

- **Interaction Social Affinity Groups (ISAGs)**: Let the set of ISAGs be the cross-product of interaction modality, action, and direction:

$$Interaction\text{-}Classes := \{link, post, photo, video\}$$
$$\times \{likes, tag, comment\}$$
$$\times \{incoming, outgoing\}$$

Then for $k \in$ *Interaction-Classes* we define

ISAG(u, k) $:= \{v | $user v has had interaction k with $u\}$

For example,

- *ISAG(u, link-like-incoming)* is the set of all users who have liked a link posted *by* user u, and

- *ISAG(u, photo-comment-outgoing)* is the set of all users whose photos user u has commented *on*.

- **Activity Social Affinity Groups (ASAGs)**: We define ASAGs based on group membership, page likes and user favourites (of which there are over 17000 distinct activities in our data set). For any one of these activities $k \in$ *Activity-Groups* we define:

ASAG(k) $:= \{v | $user v has taken part in activity $k\}$

For example,

- *ASAG(page-Beyonce)* is the set of all users who have liked *Beyonce*'s Facebook *page*, and

- *ASAG(group-Harvard Debate Club)* is the set of all users who have joined the Facebook *group* for the *Harvard Debate Club*.

---

[1] From Facebook Blog: http://www.facebook.com/blog/blog.php?post=324706977130, "Groups are the place for small group communication and for people to share their common interests and express their opinion. Groups allow people to come together around a common cause, issue or activity to organise, express objectives, discuss issues, post photos and share related content."

[2] From Facebook Blog: (http://www.facebook.com/blog/blog.php?post=324706977130 "Facebook Pages enable public figures, businesses, organisations and other entities to create an authentic and public presence on Facebook. Facebook Pages are visible to everyone on the Internet by default. Facebook users can connect with these Pages by becoming a fan and then receive their updates and interact with them."

[3] According to Facebook Blog, (https://www.facebook.com/help/232262810142682 "Facebook facilitates a wide variety of user selected favourites (Activities, Favorite Athletes, Books, Interests, Movies, Music, Sports, Favorite Teams, Television). These favourites allow a user to associate themselves with other people who share their same favourite tendencies."

2.3 Social Affinity Filtering (SAF)

With SAGs now defined, we can use them to build features for a classification-based approach to social recommendation that we term *social affinity filtering (SAF)*. In SAF, our goal is to predict $likes(u, i)$ for user u and item i. As features $X_k^{u,i}$ for this classification task, we can use the observed preferences of members of each SAG k as proxies for $likes(u, i)$. Formally, we define such features as follows:

- **Interaction Social Affinity Features (ISAFs)**: We define feature $X_k^{u,i} \in \{true, false\}$ for user u, item i and interaction $k \in \textit{Interaction-Classes}$ as

$$X_k^{u,i} := \begin{cases} true & \exists v \in ISAG(u, k) \wedge likes(v, i) = true \\ false & otherwise \end{cases}$$

In short, $X_k^{u,i}$ is *true* if any user sharing interaction k with u liked i. Here, v is implicitly limited to u's Facebook friends (with whom u may interact).

- **Activity Social Affinity Features (ASAFs)**: We define feature $X_k^{u,i} \in \{true, false\}$ for user u, item i and activity $k \in \textit{Activity-Groups}$ as

$$X_k^{u,i} := \begin{cases} true & u \in ASAG(k) \wedge \\ & \exists v \in ASAG(k) \wedge likes(v, i) = true \\ false & otherwise \end{cases}$$

In short, $X_k^{u,i}$ is *true* if both u and some other v are a member of activity k and v has liked i. Here, v may range over all Facebook users, i.e., v need not be a friend of u to share the same public activity k.

While other non-binary definitions of ISAFs and ASAFs are certainly possible (e.g., the count or fraction of members in a SAG who like the item), simple binary features provided the best performance in our experimental evaluation.

Concatenating these ISAFs and ASAFs into feature vector $\mathbf{X}(u, i) = \langle \cdots, X_k^{u,i}, \cdots \rangle$ for $k \in \textit{Interaction-Classes} \cup \textit{Activity-Groups}$ (or any subset thereof), a SAF classifier is then simply a function

$$f : \mathbf{X}(u, i) \rightarrow likes(u, i),$$

where we restrict $likes(u, i) \in \{true, false\}$.[4] Given a dataset of historical observations $D = \{\mathbf{X}(u, i) \rightarrow likes(u, i)\}$, we can *train f* using any existing classification method; in this work we consider linear classifiers trained by an SVM, logistic regression, or naïve Bayes. For *prediction*, given user u and item i, we build the feature vector $\mathbf{X}(u, i)$ and predict $likes(u, i) = f(\mathbf{X}(u, i))$ using the trained classifier f.[5]

To understand how SAF works, it helps to visualise the training data as shown in Fig 2.

---

[4] For training purposes, we omit any unobserved cases for which the class label $likes(u, i) = unknown$. At prediction time, the binary classifier must always select a class label of *true* or *false*.

[5] Since almost all classification methods provide a score (or probability) of a classification, we can also generate the top-n item recommendations for a user u by sorting items by their score.

Social Affinity Features

| | ISAF | | | | ASAF (Page) | | | | | | |
	X link-like-incoming (X_1)	X link-like-outgoing (X_2)	X video-tag-outgoing (X_{22})	X Friend liked (X_{23})	X Big Bang Theory (X_1)	X Facebook (X_2)	X Shakira (X_2)	X ANU CECS (X_{k-1})	X I Love Canberra (X_k)	Y Like / Dislike ?
u_1, i_5	1	0	0	1	0	0	1	0	0	1
u_1, i_8	0	0	1	0	0	0	0	0	0	0
u_2, i_5	1	1	0	1	0	1	0	1	1	1
u_2, i_7	0	1	1	0	1	0	0	1	0	1
u_m, i_n	0	0	1	0	0	0	1	0	0	0

Figure 2: SAF training data example: each row corresponds to a training data sample for a specific user-item pair (u, i) for which the prediction target $likes(u, i)$ is observed (last column). All other columns represent the value of ISAF or ASAF features evaluated relative to the (u, i) label of each row. All columns are binary-valued ($0 = false, 1 = true$).

2.4 SAF vs. Other Filtering Methods

While a classification approach to recommendation might evoke comparisons to *content-based filtering* (CBF) [18], we remark that CBF is not a *social* recommendation approach and unlike CBF, SAF does not require explicit user features (e.g., age, gender, location, etc.) or item descriptors (link text, link genre, etc.); in contrast, SAF uses interaction and/or activity data for social network users to define SAGs and learns the affinities between a user (ego) and the different set of alters as defined by these SAGs. Additionally, unlike state-of-the-art *social collaborative filtering* approaches [9, 19, 20, 21, 22, 23, 34], SAF does not aggregate user-user interaction and shared activity data into a single aggregate statistic, instead it uses fine-grained distinctions in this social data to define a large number of SAGs and learns which of these SAGs are informative for recommendation.

3. EVALUATION

3.1 Data Description

We built a Facebook App to collect explicit user like and dislike preferences back for links posted on Facebook (e.g., Youtube video, news or blog item, etc.) as well as detailed user interaction and activity history available through the Facebook Graph API. The data collection was performed with full permission from the user and in accordance with an approved Ethics Protocol #2011/142 from the Australian National University.

Our App requested to collect information on profiles (including activity memberships) and timelines (interactions) for the App users *and* their friends as required by Sec 2.2 and Sec 2.3. With such expressive permissions, many poten-

	App Users	Ego network of App Users
Total	119	37,872
Male	85	20,840
Female	34	17,032

Table 1: App user demographics. The *ego network* is the friend network of the App users.

App Users	Tags	Comments	Likes
Post	7,711	22,388	15,999
Link	—	7,483	6,566
Photo	28,341	10,976	8,612
Video	2,525	1,970	843
Ego network of App Users	Tags	Comments	Likes
Post	1,215,382	3,122,019	1,887,497
Link	—	891,986	995,214
Photo	9,620,708	3,431,321	2,469,859
Video	904,604	486,677	332,619

Table 2: Statistics on user *interactions*.

	App Users	Ego Network of App Users
Groups	3,469	373,608
Page Likes	10,771	825,452
Favourites	4,284	892,820

Table 3: Statistics on user *actions*, counted for *Groups, Pages* and *Favourites* over the App users and their ego network.

	Friend recommendation	Non-Friend recommendation
Like	1392	1127
Dislike	895	2111

Table 4: Dataset breakdown of prediction target $like(u, i)$ by the source of the link (Friend/Non-friend) and rating (Like=$true$, Dislike=$false$).

tial users were hesitant to install the App — after an intensive one month user drive at our University, we were able to attract 119 App users allowing us to collect activity and interaction data for a combined 37,872 users.[6]

We summarise basic statistics of the data in Tables 2–4. Table 1 presents user and friend demographics. Table 2 summarises the number of records for each item modality (row) and action (column) combination. Table 3 shows the group membership, page like and favourite counts for users.

Our App recommends three links to App users each day, which the users may optionally like or dislike. Recommended links are harvested from *both* friends' and non-friends' timelines. We display only three links per day in order to avoid rank-bias with preferences; each link could be independently rated. Table 4 shows App user link preference statistics.

All subsequent experimentats use offline batch data stored and analysed *after* a four month data collection period.

3.2 SAF Comparison

In this section, we compare novel SAF-based methods with a variety of (social) collaborative filtering baselines:

1. **Most Likely Class Constant Predictor (Const)**

2. **Nearest Neighbor (NN)** [6]

3. **Matrix Factorization (MF)** [29]

4. **Social Matchbox (SMB)** [23]

Here, Const serves as a lower bound on performance, NN and MF are two well-known state-of-the-art *non-social* collaborative filtering algorithms, and SMB is a state-of-the-art *social* collaborative filtering algorithm — a social extension of matrix factorization that encourages more homophily for pairs of users with more interactions.

Among the novel SAF methods, we analyse four different sets of social affinity features:

1. **Interaction Social Affinity Features (ISAF)**

2. **Activity-based Social Affinity Features (ASAF)** for

   (a) **Group Memberships**
   (b) **Page Likes**
   (c) **Favourites**

Furthermore, for these four classes of features, we train one of three classifier types, leading to the following classes of SAF recommenders evaluated in our experiments:

1. **Naïve Bayes (NB-ISAF, NB-ASAF)**

2. **Support Vector Machines (SVM-ISAF, SVM-ASAF)**

3. **Logistic Regression (LR-ISAF, LR-ASAF)**

In all experiments, we report average accuracy using 10-fold cross validation and provide standard error bars corresponding to 95% confidence intervals on the accuracy. Accuracy is defined as

$$accuracy = \frac{\# \ correct \ predictions}{\# \ predictions}$$

Fig 3 compares the above baselines and SAF algorithms. In all of these experiments, SAF variants performed statistically significantly better than the best baseline (SMB), except for NB-ASAF which we conjecture is due to violation of feature independence assumptions that become more pronounced as the number of features increases by two orders

---

[6] The issue of low App user uptake with such expressive App permissions underscores the importance of identifying the *minimal* set of permissions to obtain good recommendation performance — a question we address in our subsequent analysis.

Figure 3: Comparision of a simple baseline (Const), two collaborative filtering baselines (NN and MF), a social collaborative filtering baseline (SMB) and novel SAF recommenders using different feature sets (one ISAF and three ASAF sets) and classifiers (NB,LR,SVM). The best SAF-based model (LR-ASAF) — for Page likes — significantly outperforms all baselines by at least 6%. Combining all four feature sets (not shown) does not lead to improvement over Page likes features alone.

of magnitude when moving from the NB-ISAF feature set (22 features) to any of the NB-ASAF feature sets (1000's of features).

In terms of the best recommenders, we observe that LR-ASAF and SVM-ASAF perform comparably to each other and learn quite well despite the large size of this ASAF feature set. Overall, *LR-ASAF performs 6% better than the best baseline for page likes*. We combined all four features sets in a fifth experiment (not shown) and remark that it did not outperform LR-ASAF over page likes. Hence, page likes would be the most informative feature set to collect if one wanted to minimise the permissions that an App was requesting. On the other hand, we note that even if only ISAF features were available, the performance of all ISAF variants is still sufficient to statistically significantly outperform all (social) collaborative filtering baselines.

It is important to consider why ASAF outperforms ISAF. We conjecture the reasons for this are quite simple: ISAFs can only see the *friends* of user u whereas ASAFs are able to look at all users, independent of u's friends. Hence, given the relative sparsity of friend-only data in Facebook compared to the greater Facebook population (at least the population that the App collected) and also the relative number of interaction SAGs compared to activity SAGs, ASAFs appear to draw on a much larger set of SAGs that in turn draw on a much larger user population. Among ASAF activities, page likes are the most predictive followed by group membership and favourites. This reinforces our conjecture that data sparsity can hurt SAF since we note from Table 3 that page likes are more prevalent than groups and favourites.

Comparing SAF to the state-of-the-art in social collaborative filtering as represented by Social Matchbox (SMB) [23], we observe that SAF consistently outperforms it. We note that the key difference of SAF vs. SMB is that SAF exploits the predictiveness of fine-grained interactions and activities — it breaks them down into small subgroups, whereas most social collaborative filtering approaches [9, 19, 20, 21, 22, 23, 34] instead collapse the diverse set of interactions into aggregate statistics such as the number of interactions between user u_1 and user u_2, regardless of whether u_1 tagged u_2 in a photo or u_1 liked a photo on u_2's wall. Clearly there is a great deal of benefit to be derived from these fine-grained interactions and it suggests that one might rethink existing social collaborative filtering approaches that do aggregate interaction and activity information into aggregate statistics.

On two final notes, we remark that SAF yields a computational and optimization advantage over (social) collaborative filtering in that it is straightforward and efficient to find a globally optimal classifier with respect to certain training criteria (e.g., optimising log loss in logistic regression or hinge loss in SVMs) unlike (social) collaborative filtering approaches that generally rely on computationally expensive nearest neighbor or matrix factorization techniques that lack training optimality guarantees. Further, we also note that SAF inherently scales to a large number of users and generalizes to completely new users without suffering from the cold-start problem: this is simply because nothing SAF learns is user-dependent, it learns to weight SAGs independent of individual users.

Given the clearly demonstrated benefits of SAF, we now

proceed in the next two sections to analyse the two primary types of SAG features (interactions and activities) to better understand characteristics of both informative and uninformative SAGs in each context and the social phenomena that may be responsible for these characteristics.

3.3 Cold-start Analysis

Collaborative filtering algorithms suffer from the user cold-start problem, where no historical information about user is available. One of the advantage of SAF is that the social affinity features for a user-item are defined in terms of the user's affinity groups(ASAG/ISAG). Hence, the SAF learning is not highly dependent of individual user's data.

For user cold-start analysis, we create 10 fold train-test set where users in train and test set are mutually exclusive and each test set consists of 10% of the total users. We hold out 30% of each user's data from the test set and train two SAF predictors namely: *cold-start* and *non cold-start* predictor. We train cold-start predictor using training set and non cold-start predictor with additional held out data from test set. Finally, we evaluate the performance of cold-start and non cold-start predictor on remaining 70% of test data. In fig 4 we clearly see that the accuracy[7] of SAF predictor for cold-start is significantly better than constant predictor. Furthermore, the accuracy of cold-start predictor is comparable to non cold-start predictor, which indicates that SAF performs quite well for cold-start users where most of the existing methods fail. Hence, unlike standard (social) collaborative filtering techniques, SAF is robust to user cold-start problem.

3.4 Interaction Analysis

In this section we analyse the informativeness of Interaction Social Affinity Features (ISAFs), namely user interactions according to their modality, type, and direction, as described in Sec 2.

A general method for measuring the amount of information that a feature $X_k^{u,i}$ provides w.r.t. predicting a user preference $likes(u,i)$ (in this case, just $true$ or $false$) is to calculate its conditional entropy:

$$H(likes(u,i)|X_k^{u,i} = true)$$
$$= - \sum_{y \in (true,false)} p(likes(u,i) = y|X_k^{u,i} = true)$$
$$\cdot \ln(p(likes(u,i) = y|X_k^{u,i} = true))$$

In general, a lower conditional entropy indicates a more informative feature. Here we measure the conditional entropy $H(likes(u,i)|X_k^{u,i} = true)$ rather than mutual information $I(likes(u,i); X_k^{u,i})$, as we found that mutual information is highly correlated with (and dominated by) the frequency of the feature $X_k^{u,i} = true$ in the dataset.

---

[7] The slight decrease in accuracy for non cold-start case compared to fig 3 is due to the fact that the training set consists of only 30% of test set users data and test set is biased to small set of test users.

Table 5: Conditional entropy of various interactions (lower conditional entropies are more informative).

Modality (X)	$H(Y|X = true)$
video	0.850
link	0.915
post	0.918
photo	0.926

Action Type (X)	$H(Y|X = true)$
tags	0.920
comments	0.921
likes	0.924

Direction (X)	$H(Y|X = true)$
outgoing	0.928
incoming	0.935

Modality-Direction (X)	$H(Y|X = true)$
tags-outgoing	0.885
likes-outgoing	0.885
tags-incoming	0.900
likes-incoming	0.902
comments-outgoing	0.908
comments-incoming	0.912

| Action-Direction (X)	$H(Y|X = true)$
photo-outgoing	0.857
video-outgoing	0.863
link-outgoing	0.895
link-incoming	0.896
post-incoming	0.902
post-outgoing	0.906
video-incoming	0.915
photo-incoming	0.921

Figure 4: Comparision of the SAF for user cold-start and non cold-start cases. Accuracy of SAF predictor evaluated on cold-start users outperforms constant predictor baseline and is comparable to the non cold-start SAF predictor.


Figure 5: Conditional Entropy of modalities/activities for incoming/outgoing interactions vs. item liked by at least k friends. Increasing k generally has a stronger influence on informativeness than other features of interaction SAGs (modality, action, or direction), with the exception of the outgoing-modality.

First we analyse various interactions to understand what interactions define SAGs with a high affinity for a user u's preferences. To this end, we make a few observations from the conditional entropy analysis of Table 5:

- Interaction on *videos* seem to have a stronger preferential affinity than other modalities such as links, posts and photos. This could be due to the fact that video viewing is time-consuming and users inherently only watch the videos of those whose preferences they often share.

- Tagging has a slightly better conditional entropy than commenting and liking, potentially because tagging a user often indicates a direct social interaction (appearing in a photo or video together) that provides evidence of affinity between two users.

- A user is more likely to share preferences with someone who she initiates the interaction with (outgoing) vs. with someone who initiates the interaction with her (incoming). As an extreme instance of this, we note that while outgoing photo and video interactions are *most* informative in the last table of Table 5, it appears that incoming photo and video interactions are *least* informative.

In Fig 5 we plot the conditional entropy of modality and action for incoming/outgoing interactions constrained to links liked by at least k friends in the SAG (measuring the implicit or explicit exposure of a user to their friends preferences via a SAG). Fig 5 reiterates many observations above for various values of fixed k. In addition, we note that preference affinity with a SAG increases as more people in the SAG like the item. While incoming interactions were not as predictive as outgoing interactions for the same k, we note that higher k for an incoming interaction *can be more predictive* than lower k for an outgoing interaction. Similar principles hold for modality and action vs. k — a larger k is generally more predictive than the individual variation among modality and action, the one exception being the modality-outgoing analysis. Overall, these observations suggest a large cumulative number of friend preferences in an interaction SAG can be more predictive than other features of the interaction SAG. Further investigation is needed to pinpoint whether or not there are diminishing returns on repeated exposures [27, 31] on k, and how this could be leveraged into future feature engineering in SAF-based recommender system design.

3.5 Activity Analysis

Now we analyse the informativeness of Activity Social Affinity Features (ASAFs) by looking at the correlation between the size and type of groups, pages and favourites.

Fig 6 shows the relationship between both the conditional entropy and logistic regression weights vs. the size of activity groups. Here the size of a *group*, *page* and *favourite* is the number of total users in the activity group. For *pages* and *favorites* this is the total number of Facebook users, whether or

not they are in the App users' ego network, while for *groups* only the number of users in the App users' ego network is visible to our app. Both scatter plots shows that the activity groups of small size can be highly predictive (low conditional entropy or weights that deviate extremely from zero) whereas large groups are rarely predictive.

In Fig 7 we plot the average conditional entropy of the top 10% of features cumulative up to the size of the activity group given on the x-axis; this allows us to determine the marginal contribution of larger groups to the average conditional entropy as larger groups are incrementally added in. This graph distinctly shows that the small sizes of groups, pages and favourites have low average conditional entropy that transitions sharply to a higher average once a size threshold has been met. From Fig 7 we can infer that the group sizes up to 50 and page/favourite sizes up to 10^5 are most predictive.

We also analyse predictiveness of favourites by categories in Fig 8, where the favorite category labels are obtained from the Facebook API. We can see that contents in the "longtail", i.e., having a large number of occurrences far from the most popular choices, tend to have some of the most predictive individual affinities. Examples of these include music, books, movies. On the contrary, generic affinities (e.g. interests) and those with a smaller number of choices (e.g. sports or fav-teams) tend to be less predictive since they represent less specialised interests than the long tail of music, book, or movie preferences.

These observations of Fig 8 are also reiterated by the examples provided in Table 6 where uninformative favourites tend to have a broad appeal whereas informative favourites generally appear much more specialised. This also reinforces the point that not all SAGs are predictive, but some are very predictive and it is important to learn which SAGs are informative rather than naïvly aggregate their content, where on average, the features are clearly not informative.

In Fig 9 we analyse the relationship between accuracy and number of active features i.e. features that are true. We can see that accuracy increase as number of active features increases but then starts to decrease sharply. This is due to the fact that items with large number of active features are likely to be general items, liked by wide variety of users, which makes it hard for SAF to make correct prediction. In Fig 10 we can see that , in general, accuracy of SAF increases as number of group memebership, page likes and favourites increase. This indicates that SAF can make better recommendation to those users who actively express their preferences in a social network via group memebership, page likes and favourites.

4. RELATED WORK

This work relates to many others in inferring user preferences on social and information networks. We structure the discussion into three parts: the first is concerned with the nature and observations on user traits, interactions and dif-
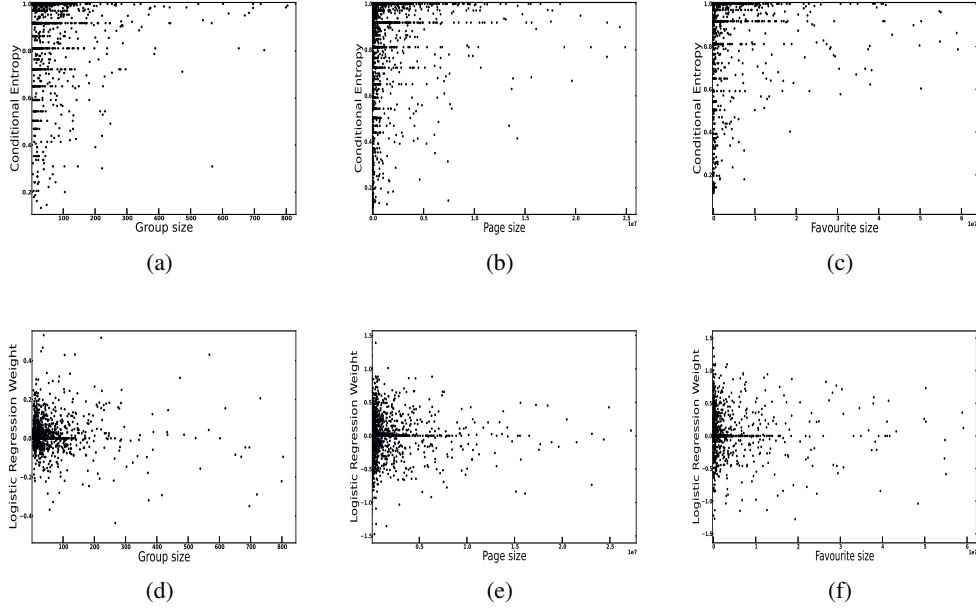
Figure 6: Conditional entropy vs size (a-c); logistic regression feature weights vs size (d-f). In (a-c) we observe that the large membership ASAGs are rarely informative while the most informative SAGs tend to have low memberships. Similarly in (d-f) we see that the most predictive features with the most extreme weights are concentrated toward small ASAGs.


Figure 7: Average conditional entropy of top 10% groups, pages and favourite features *cumulative* over the size. Here we see that as we add in larger membership ASAGs, the average informativeness decreases substantially (entropy increases).


Figure 8: Conditional entropy for top 1000 favourites breakdown by categories. While ASAG categories with many options like music are not informative on average, we see that some of the most informative ASAGs are music. This reiterates the point that it is crucial to *learn* which ASAGs (or ISAGs) are informative rather than aggregating average information.

10

Figure 9: Accuracy increases as the number of active features increases, but then, after reaching a certain limit, it starts to decrease sharply.


Figure 10: Accuracy of the SAF increases as user becomes more active in social network by joining more groups/pages/favourites.

Median Informative Favourites by Category				
Books	**Movies**	**Music**	**Television**	**Interests**
Harry Potter series	Forrest Gump	John Lennon	Futurama	Travel
A Song of Ice and Fire	Pretty Woman	U2	Star Trek	Music
Discworld	Napoleon Dynamite	AC/DC	The Trap Door	Literature
Hitchhiker's Guide To The Galaxy	Harry Potter	The Smashing Pumpkins	Drawn Together	Painting
The Hobbit	Toy Story 3	Gotye	Sherlock(Official)	Running
The Magician's Guild	The Godfather	The Rolling Stones	Hitchhiker's Guide to the Galaxy	Sports
Ranger's Apprentice	Mulan	All Axess	Buffy The Vampire Slayer	Films
Cosmos	How to Train Your Dragon	Steve Aoki	South Park	Genetics
Foundation and Earth	The Princess Bride	Rihanna	24	Travelling
Deception Point	Watchmen	Billy Joel	The Daily Show	Internet

Most Informative Favourites by Category				
Books	**Movies**	**Music**	**Television**	**Interests**
Calvin and Hobbes	Billy Madison	Avascular Necrosis	Metalocalypse	Computers
Tomorrow when the War Began	Team America: World Police	Tortured	Beast Wars	Texas HoldEm
I really like ceilings	Pan's Labyrinth	Elysian	Hey Arnold!	Programming
Angels and demons	Pirates of the Caribbean	Anno Domini	Sherlock	Economics
Magician	Aladdin	Darker Half	Hey It's Saturday	Martial arts
Digital Fortress	Starship Troopers	Hellbringer	Neil Buchanan and Art Attack!	Graphic design
The Bible	Happy Gilmore	Johnny Roadkill	Breaking Bad	Cooking
Interview with the Vampire	Timon and Pumbaa	Aeon of Horus	Red vs. Blue	Klingon language
The Discworld Series	Ferris Buellers Day Off	Katabasis	Stargate Universe	Politics
The Da Vinci Code	Peter Griffin	Bane Of Isildur	Chaser's War on Everything	Science

Table 6: (top) Examples of 10 items per Favourite category near the *median* conditional entropy (*median informativeness*). (bottom) Examples of top 10 items with the lowest conditional entropy (*most informative*). A general trend is that more informative favourite category ASAGs tend to be more specialised in appeal, e.g. "Avascular Necrosis" is an informative music group favourite — its members tend to share common preferences — while "John Lennon" and "U2" have a broader audience with more diverse preferences. Interestingly, "Sherlock" appears in both most and median informative table but the median informative is an official page with wide range of fans, whereas the most informative is a duplicate fan page with few number of fans.

fusion mechanisms; the second is concerned with correlating these user traits and interactions to user preferences and interests; the third is concerned with methods that uses these observations for predicting user interest or recommending content on social networks.

The first group of related work studies the nature of user profile, interactions, and diffusion. Profile information and demographics is correlated with user behavior patterns. Chang *et al* [8] showed that the tendency to initiate a Facebook friendship differs quite widely across ethnic groups, while Backstorm *et al* [3] have additionally showed that female and male users have opposite tendencies for dispersing attention for within-gender and across-gender communication. Two particular measurement studies on Facebook attention [3, 33] have inspired our work. Although the average number of friends for a Facebook user is close to the human psychological limit, known as the Dunbar number [14], the findings concur that a user's attention (i.e., interactions) are divided among a much smaller subset of Facebook friends. [3] studied two types of attention: communication interaction and viewing attention (e.g. looking at profiles or photos). Users' communication attention is focused on small numbers of friends, but viewing attention is dispersed across all friends. This finding supports our approach of looking at many types of user interactions across all of a user's contact network, as a user's interest is driven by where he or she focuses attention on.

The mechanisms of diffusion invites interesting mathematical and empirical investigations. The Galton-Watson epidemics model suits the basic setup of social message diffusion, and can explain real-world information cascade such as email chain-letters when adjusted with selection bias [12]. For social diffusions in a one-to-many setting, however, the epidemics model has been less accurate. Ver Steeg *et al* [31] found that online message cascades (on Digg social reader) are often smaller than prescribed by the epidemics model, seemingly due to the diminishing returns of repeated exposure. Romero *et al* [27], in an independent study, confirmed the effect of diminishing returns with Twitter hashtag cascades, and further found that cascade dynamics differ across broad topic categories such as politics, culture, or sports. Our observations on user preference on items liked by a number of Facebook friends suggest large cumulative number of friend interactions is more predictive, further investigation is needed to pinpoint the effect of diminishing returns on repeated exposures.

The nature of social diffusion seem to be not only democratic [2, 4], but also broadening for users [5]. While influential users are important for cascade generation [4], large active groups of users are needed to contribute for the cas-

cade to sustain [2]. Moreover, word-of-mouth diffusion can only be harnessed reliably by targeting large numbers of potential influencers, confirmed by observations on Twitter [4] and online ads [32]. In a study facilitated by A/B testing on Facebook links, [5] found that while people are more likely to share the information they were exposed to by their strong ties than by their weak ties, the bulk of information we consume and share comes from people with different perspectives (weak ties). Our Facebook App is intended to bridge this gap between insights from these observations and predicting user actions.

The second group of related work tries to correlate from user interactions to preferences and tie strength. Saez-Trumper et al [28] found that incoming and outgoing actives are highly correlated on broadcast platforms such as Facebook and Twitter, and such correlation does not hold in one-to-one mode of communication such as email. Multiple studies have found that online interactions tend to correlate more with interests than with user profile. Singla *et al* [30] found that user who frequently interact (via MSN chat) tend to share (web search) interests. Anderson *et al* [1] concluded that the level of user activities correlate with the positive ratings that they give each other, i.e., it is less about what they say (content of posts) but more about who they interacted with. Such findings echo those by Brandtzag [7] that real-world interactions (e.g., appearing in the same photo) further strengthens friendship on Facebook, while virtual interactions reveal interests. Furthermore, ratings of real-world friendship strength and trust [10] seems to be better predicted from the intimacy, intensity, and duration of interactions, than from social distance and network structure. Our work is not only inspired by these observations, we also quantify the strength of correlations of user interest with a large variety of user affinities – namely, activities, and group preferences in different categories.

The last group of related work is concerned with using social network and behavior information for recommendation. Matrix factorization is one of the prevailing approaches for recommender systems [17, 21]. Recent advances include extending matrix factorization to user social relation in regularization [22, 19], to take into account multiple relations [26, 16], and to model social context [15]. In particular, there are different designs for using social information to regularize objective functions [34], a trust ensemble [20], a low-rank factorization of the social interactions matrix [21], or social-spectral regularization that takes into account user and item features [23]. These systems have shown very promising performance across a range of problems, but their all collapse social affinity (fine-grained interactions and group affinity) into one or a very low-dimensional representation. The point of departure of this work is to explore the rich affinity structure, we compared a recent matrix factorization approach [23] and found SAF with simple classifiers outperform state-of-the-art.

Additional work on predicting user actions join multiple social networks and explores logical representation of user actions. Nori *et al* [24] examines predictability of user actions on Twitter from actions in Twitter and Del.icio.us. The study uses both linear regression and a bipartite graph model that outperformed state-of-the-art models. Gomes *et al* [13] derived rules for Facebook interactions using a psychology-inspired formal symbolic language. Our affinity definition is based on direct interactions within a users' ego network, this is complementary to a recent alternative [25] that uses number of paths between two users encodes the resilience of network structure, as it was recently found [11] that the vast majority of information diffusion happens within one step from the source node. These work are most closely related to ours, yet none has examined such a diverse set of user actions in the same context: one-on-one interactions (e.g. commenting), broadcast (e.g. posting, sharing), and co-preference (e.g. likes).

In summary, our study is motivated by overall utility of diverse and fine-grained user interactions. To the best of our knowledge, this is the first work that look at 10,000+ different types of social affinity. We show that rich affinity features outperform state-of-the-art recommendation approaches, and our observations confirm the effect of diminishing returns on repeated exposure, we observe that contents in the *long tail* tend to be more predictive, and quantified the correlation of a large variety of affinity traits with user preferences.

5. CONCLUSIONS

We proposed Social Affinity Filtering (SAF) as a novel method for social recommendation that analyses a user's fine-grained interactions and activities to learn which subset of social groups have the highest affinity with a user's preferences. We evaluated SAF on a dataset collected from a Facebook App, showing that SAF yields 6% absolute improvement in accuracy over state-of-the-art (social) recommendation engines just using knowledge of users' page likes. This is an important result given that SAF is built on standard supervised linear classification techniques (which support strong training guarantees and fast learning algorithms) and would be fairly robust in an online setting in contrast to some more complex matrix factorization optimisation approaches often proposed for (social) collaborative filtering.

In addition to the strong social recommendation accuracy improvements offered by SAF, we quantified the relative importance of interaction and activity groups for recommendation and we analysed what properties made some social affinity groups more informative than others. Among many insights, our results show that video interactions are more predictive than other modalities and outgoing interactions are more predictive than incoming ones, but both can be superceded by the number of preferences expressed by a group. Furthermore, for activities, we showed that smaller social groups are more predictive than larger ones and long-tailed categories with many specialised choices (dynamically increasing over time) tend to contain some of the most pre-

dictive affinity groups. Knowing what subset of features are informative allows one to design an effective recommendation tool that requires minimal permissions from a user — a key property for general user uptake.

Future directions of research can examine the nature of social groups via additional metrics — e.g., the social network within members of the group, or activity level of the group. Other work might explore the feature engineering to better incorporate preference frequency, or even combinations of SAF with orthogonal (social) collaborative filtering approaches like nearest neighbor or even matrix factorization.

6. ACKNOWLEDGEMENTS

7. REFERENCES

[1] A. Anderson, D. Huttenlocher, J. Kleinberg, and J. Leskovec. Effects of User Similarity in Social Media. In *WSDM*, 2012.

[2] S. Asur, B. A. Huberman, G. Szabo, and C. Wang. Trends in social media: Persistence and decay. In *ICWSM*, 2011.

[3] L. Backstrom, E. Bakshy, J. Kleinberg, T. Lento, and I. Rosenn. Center of attention: How facebook users allocate attention across friends. In *ICWSM*, 2011.

[4] E. Bakshy, J. M. Hofman, W. A. Mason, and D. J. Watts. Everyone's an influencer: quantifying influence on twitter. WSDM '11, 2011.

[5] E. Bakshy, I. Rosenn, C. Marlow, and L. Adamic. The role of social networks in information diffusion. *Facebook report, http://www.scribd.com/facebook*, 2012.

[6] R. M. Bell and Y. Koren. Scalable collaborative filtering with jointly derived neighborhood interpolation weights. In *ICDM-07*, 2007.

[7] P. B. Brandtzg and O. Nov. Facebook use and social capital — a longitudinal study. ICWSM'11, 2011.

[8] J. Chang, I. Rosenn, L. Backstrom, and C. Marlow. epluribus : Ethnicity on social networks. In *ICWSM '10*, pages 18–25, 2010.

[9] P. Cui, F. Wang, S. Liu, M. Ou, and S. Yang. Who should share what? item-level social influence prediction for users and posts ranking. In *SIGIR*, 2011.

[10] E. Gilbert and K. Karahalios. Predicting tie strength with social media. In *Proc. CHI*. ACM, 2009.

[11] S. Goel, D. J. Watts, and D. G. Goldstein. The structure of online diffusion networks. In *EC*, EC '12, pages 623–638, New York, NY, USA, 2012. ACM.

[12] B. Golub and M. O. Jackson. Using selection bias to explain the observed structure of internet diffusions. *Proc. Nat. Academy Sci.*, 107(24), 2010.

[13] A. Gomes and M. da Graca C Pimentel. Social interactions representation as users behavioral contingencies and evaluation in social networks. In *ICSC*. IEEE, 2011.

[14] R. Hill and R. Dunbar. Social network size in humans. *Human Nature*, 14(1):53–72, 2003.

[15] M. Jiang, P. Cui, R. Liu, Q. Yang, F. Wang, W. Zhu, and S. Yang. Social contextual recommendation. CIKM '12, pages 45–54, 2012.

[16] M. Jiang, P. Cui, F. Wang, Q. Yang, W. Zhu, and S. Yang. Social recommendation across multiple relational domains. CIKM '12, 2012.

[17] Y. Koren, R. Bell, and C. Volinsky. Matrix factorization techniques for recommender systems. *Computer*, 42, 2009.

[18] K. Lang. NewsWeeder: Learning to filter netnews. In *ICML-95*, 1995.

[19] W.-J. Li and D.-Y. Yeung. Relation regularized matrix factorization. In *IJCAI-09*, 2009.

[20] Ma, King, and Lyu. Learning to recommend with social trust ensemble. In *SIGIR-09*, 2009.

[21] H. Ma, H. Yang, M. R. Lyu, and I. King. Sorec: Social recommendation using probabilistic matrix factorization. In *CIKM-08*, 2008.

[22] H. Ma, D. Zhou, C. Liu, M. R. Lyu, and I. King. Recommender systems with social regularization. In *WSDM-11*, 2011.

[23] J. Noel, S. Sanner, K.-N. Tran, P. Christen, L. Xie, E. V. Bonilla, E. Abbasnejad, and N. Della Penna. New objective functions for social collaborative filtering. In *WWW*, pages 859–868, New York, NY, USA, 2012. ACM.

[24] N. Nori, D. Bollegala, and M. Ishizuka. Exploiting user interest on social media for aggregating diverse data and predicting interest. ICWSM '11, 2011.

[25] R. Panigrahy, M. Najork, and Y. Xie. How user behavior is related to social affinity. WSDM '12, 2012.

[26] S. Rendle, L. B. Marinho, A. Nanopoulos, and L. Schmidt-Thieme. Learning optimal ranking with tensor factorization for tag recommendation. In *KDD-09*, 2009.

[27] D. M. Romero, B. Meeder, and J. Kleinberg. Differences in the mechanics of information diffusion across topics: idioms, political hashtags, and complex

contagion on twitter. WWW '11, 2011.

[28] D. Saez-Trumper, D. Nettleton, and R. Baeza-Yates. High correlation between incoming and outgoing activity: A distinctive property of online social networks? In *ICWSM*, ICWSM '11, 2011.

[29] R. Salakhutdinov and A. Mnih. Probabilistic matrix factorization. In *NIPS 20*, 2008.

[30] P. Singla and M. Richardson. Yes, there is a correlation: - from social networks to personal behavior on the web. WWW '08, 2008.

[31] G. Ver Steeg, R. Ghosh, and K. Lerman. What stops social epidemics? ICWSM '11, 2011.

[32] D. J. Watts and P. S. Dodds. Influentials, networks, and public opinion formation. *Journal of Consumer Research*, 2007.

[33] C. Wilson, B. Boe, A. Sala, K. Puttaswamy, and B. Zhao. User interactions in social networks and their implications. EuroSys'09, 2009.

[34] S.-H. Yang, B. Long, A. Smola, N. Sadagopan, Z. Zheng, and H. Zha. Like like alike: Joint friendship and interest propagation in social networks. In *WWW-11*, 2011.